

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb

# Disable scientific notation for large numbers
pd.options.display.float_format = '{:.0f}'.format

# Setting display options for Pandas to show three decimal places for floati
pd.set_option('display.float_format', lambda x: '%.2f' % x)
```

Data Loading

```
In [2]: # import data
sales_df = pd.read_csv('/content/drive/MyDrive/Global_Superstore.csv', encoc
```

Data Exploration

```
In [3]: sales_df.info() # Display information about the DataFrame, including data ty
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 51290 entries, 0 to 51289
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Row ID                 51290 non-null  int64
1   Order ID               51290 non-null  object
2   Order Date             51290 non-null  object
3   Ship Date              51290 non-null  object
4   Ship Mode              51290 non-null  object
5   Customer ID            51290 non-null  object
6   Customer Name          51290 non-null  object
7   Segment                51290 non-null  object
8   City                   51290 non-null  object
9   State                  51290 non-null  object
10  Country                 51290 non-null  object
11  Postal Code             9994 non-null   float64
12  Market                  51290 non-null  object
13  Region                  51290 non-null  object
14  Product ID              51290 non-null  object
15  Category                51290 non-null  object
16  Sub-Category            51290 non-null  object
17  Product Name            51290 non-null  object
18  Sales                   51290 non-null  float64
19  Quantity                51290 non-null  int64
20  Discount                51290 non-null  float64
21  Profit                  51290 non-null  float64
22  Shipping Cost           51290 non-null  float64
23  Order Priority           51290 non-null  object
dtypes: float64(5), int64(2), object(17)
memory usage: 9.4+ MB
```

```
In [4]: sales_df.head() # Display top 5 records
```

```
Out[4]:
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	
0	32298	CA-2012-124891	31-07-2012	31-07-2012	Same Day	RH-19495	Rick Hansen	Consumer	Ne
1	26341	IN-2013-77878	05-02-2013	07-02-2013	Second Class	JR-16210	Justin Ritter	Corporate	Wollo
2	25330	IN-2013-71249	17-10-2013	18-10-2013	First Class	CR-12730	Craig Reiter	Consumer	Br
3	13524	ES-2013-1579342	28-01-2013	30-01-2013	First Class	KM-16375	Katherine Murray	Home Office	
4	47221	SG-2013-4320	05-11-2013	06-11-2013	Same Day	RH-9495	Rick Hansen	Consumer	

5 rows x 24 columns

```
In [5]: sales_df[['Sales', 'Quantity', 'Discount', 'Profit', 'Shipping Cost']].describe()
```

```
Out[5]:
```

	Sales	Quantity	Discount	Profit	Shipping Cost
count	51290.00	51290.00	51290.00	51290.00	51290.00
mean	246.49	3.48	0.14	28.61	26.38
std	487.57	2.28	0.21	174.34	57.30
min	0.44	1.00	0.00	-6599.98	0.00
25%	30.76	2.00	0.00	0.00	2.61
50%	85.05	3.00	0.00	9.24	7.79
75%	251.05	5.00	0.20	36.81	24.45
max	22638.48	14.00	0.85	8399.98	933.57

```
In [6]: sales_df.isna().sum() # Find sum of missing values
```

Out[6]:

	0
Row ID	0
Order ID	0
Order Date	0
Ship Date	0
Ship Mode	0
Customer ID	0
Customer Name	0
Segment	0
City	0
State	0
Country	0
Postal Code	41296
Market	0
Region	0
Product ID	0
Category	0
Sub-Category	0
Product Name	0
Sales	0
Quantity	0
Discount	0
Profit	0
Shipping Cost	0
Order Priority	0

dtype: int64

```
In [7]: sales_df.fillna(0,inplace=True) # Replace missing values with 0
```

```
In [8]: sales_df.isna().sum() # Find sum of missing values after cleaning
```

```
Out[8]:
```

	0
Row ID	0
Order ID	0
Order Date	0
Ship Date	0
Ship Mode	0
Customer ID	0
Customer Name	0
Segment	0
City	0
State	0
Country	0
Postal Code	0
Market	0
Region	0
Product ID	0
Category	0
Sub-Category	0
Product Name	0
Sales	0
Quantity	0
Discount	0
Profit	0
Shipping Cost	0
Order Priority	0

dtype: int64

```
In [9]: print(sales_df.duplicated().sum()) # Find sum of duplicate records
```

0

Since, sum of duplicated values is zero, so there is no need to drop duplicates.

```
In [10]: # Convert date column to datetime
sales_df['Order Date'] = pd.to_datetime(sales_df['Order Date'])
sales_df['Ship Date'] = pd.to_datetime(sales_df['Ship Date'])
```

```

/tmp/ipython-input-10-1068123818.py:2: UserWarning: Parsing dates in %d-%m-%Y format when dayfirst=False (the default) was specified. Pass `dayfirst=True` or specify a format to silence this warning.
  sales_df['Order Date'] = pd.to_datetime(sales_df['Order Date'])
/tmp/ipython-input-10-1068123818.py:3: UserWarning: Parsing dates in %d-%m-%Y format when dayfirst=False (the default) was specified. Pass `dayfirst=True` or specify a format to silence this warning.
  sales_df['Ship Date'] = pd.to_datetime(sales_df['Ship Date'])

```

In [11]: `sales_df.head()`

Out[11]:

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	
0	32298	CA-2012-124891	2012-07-31	2012-07-31	Same Day	RH-19495	Rick Hansen	Consumer	Ne
1	26341	IN-2013-77878	2013-02-05	2013-02-07	Second Class	JR-16210	Justin Ritter	Corporate	Wollc
2	25330	IN-2013-71249	2013-10-17	2013-10-18	First Class	CR-12730	Craig Reiter	Consumer	Br
3	13524	ES-2013-1579342	2013-01-28	2013-01-30	First Class	KM-16375	Katherine Murray	Home Office	
4	47221	SG-2013-4320	2013-11-05	2013-11-06	Same Day	RH-9495	Rick Hansen	Consumer	

5 rows × 24 columns



In [12]: `# Save cleaned data`
`sales_df.to_csv('/content/drive/MyDrive/Global_Superstore_final.csv', index=`


Data Visualization

In [13]: `# Explore distributions of numerical data`
`# List of numerical columns to be explored`
`num_cols = ['Sales', 'Profit', 'Discount']`

`# Create histograms for each numerical column`
`for col in num_cols:`

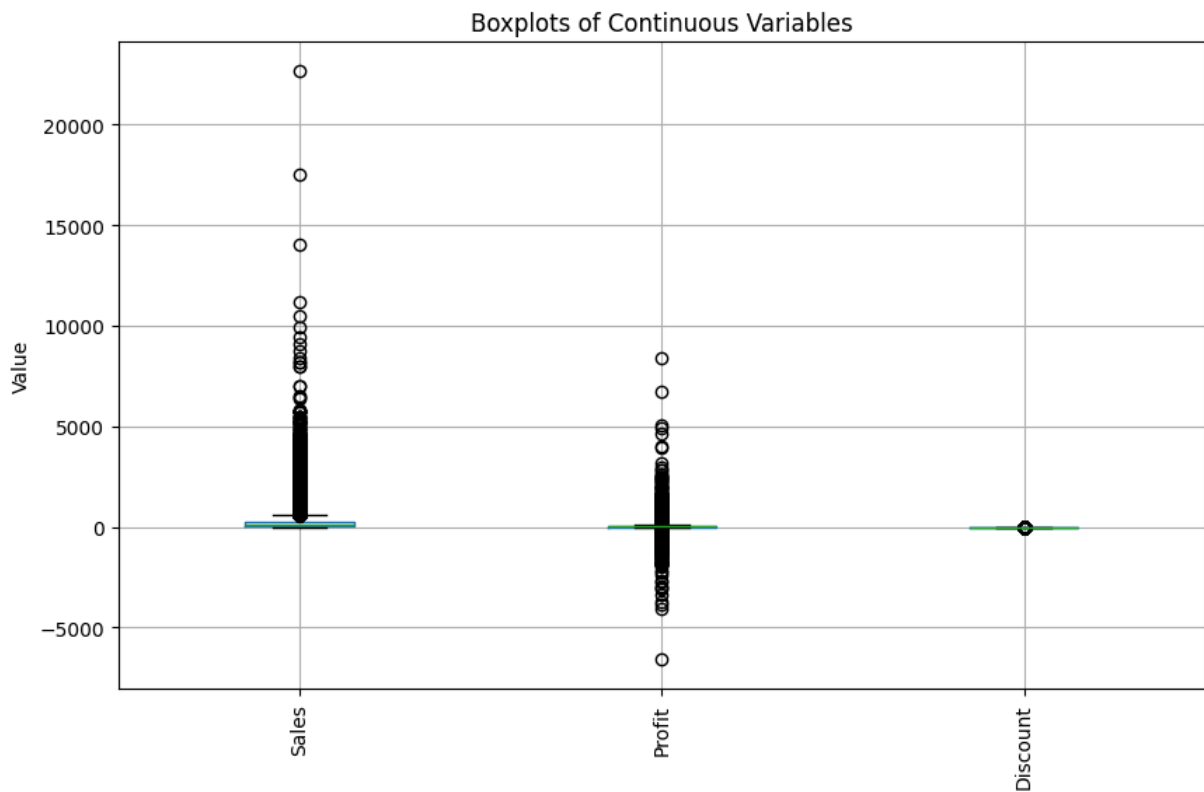
```
plt.figure(figsize=(8, 5))
sales_df[col].hist(bins=20, edgecolor='black')
plt.title(f'Distribution of {col}')
plt.xlabel(col)
plt.ylabel('Frequency')
plt.grid(False)
plt.show()
```





```
In [14]: # Identify outliers in continuous variables
# List of continuous (numerical) columns
continuous_cols = ['Sales', 'Profit', 'Discount']


# Create boxplots for each continuous variable
plt.figure(figsize=(10, 6))
sales_df[continuous_cols].boxplot()
plt.title('Boxplots of Continuous Variables')
plt.ylabel('Value')
plt.xticks(rotation=90)
plt.grid(True)
plt.show()
```



```
In [15]: # Visualize correlations and relationships between features
# Select numerical columns for correlation analysis
num_cols = ['Sales', 'Profit', 'Discount', 'Quantity']


# Calculate the correlation matrix
corr_matrix = sales_df[num_cols].corr()

# Create the heatmap
plt.figure(figsize=(8, 6))
sb.heatmap(corr_matrix, annot=True, cmap='coolwarm', center=0, fmt=".2f", li
plt.title('Feature Correlation Heatmap')
plt.show()
```





```
In [16]: # Remove outliers using IQR method
#Q1 = sales_df.quantile(0.25)
#Q3 = sales_df.quantile(0.75)
#IQR = Q3 - Q1
# Filter dataframe to keep only rows within 1.5*IQR range for all numerical
#sales_df_clean = sales_df[~((sales_df < (Q1 - 1.5 * IQR)) | (sales_df > (Q3 + 1.5 * IQR)))]
```

```
In [17]: # Observe trends in Sales over time.
plt.figure(figsize=(12, 6))
#sales_df.groupby('Order Date')['Sales'].sum().plot(kind='line', marker='o')
trend_data = sales_df.groupby(pd.Grouper(key='Order Date', freq='ME'))['Sales'].sum()
sb.lineplot(x=trend_data.index, y=trend_data.values)
plt.title('Monthly Sales Trend')
plt.xlabel('Order Date')
plt.ylabel('Total Sales')
plt.grid(True)
plt.show()
```




```
In [25]: # Relationship between Profit and Discount
plt.figure(figsize=(8, 6))
plt.scatter(sales_df['Discount'], sales_df['Profit'], alpha=0.5)
plt.xlabel('Discount')
plt.ylabel('Profit')
plt.title('Profit vs Discount')
plt.grid(True)
plt.show()
```




Visualize sales distribution by Region and Category

```
In [19]: # Plot pie chart of sales distribution by Region
region_sales = sales_df.groupby('Region')['Sales'].sum()
plt.figure(figsize=(8, 8))
plt.pie(region_sales, labels=region_sales.index, autopct='%1.1f%%', startangle=90)
plt.title('Sales Distribution by Region')
plt.axis('equal')
plt.show()
```



```
In [20]: # Plot pie chart of sales distribution by Category
category_sales = sales_df.groupby('Category')['Sales'].sum()
plt.figure(figsize=(8, 8))
plt.pie(category_sales, labels=category_sales.index, autopct='%1.1f%%', startangle=90)
plt.title('Sales Distribution by Category')
plt.axis('equal')
plt.show()
```



Train a Linear Regression Model and evaluate model performance

```
In [21]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [22]: # Feature Engineering
X = sales_df[['Profit', 'Discount']]
y = sales_df['Sales']

# Split data (training data 80% test data 20%)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train model
model = LinearRegression()
model.fit(X_train, y_train)

# Generate predictions
y_pred = model.predict(X_test)
```

```
# Calculate evaluation metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print model performance
print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared: {r2:.2f}")
```

Mean Squared Error: 187787.80

R-squared: 0.16

This notebook was converted with convert.ploomber.io